



# Synql: A CRDT-based Approach for Replicated Relational Databases with Integrity Constraints

Victorien Elvinger, Claudia-Lavinia Ignat, Habibatu Ba  
Inria Nancy, France



June 2024

# Adding collaborative features to applications




- several persons **modify together** a shared content
  - located at **different places**
  - **simultaneous** modifications or at **distinct time**
- adding collaborative features to applications is hard
  - **sequential** → **concurrent** modifications
  - **offline support**

# Adding collaborative features to applications



- replicate the application  $\implies$  dedicated development

# Adding collaborative features to applications




- replicate the application  $\implies$  dedicated development
- replicate the application data<sup>a</sup>

---

<sup>a</sup>Kleppmann et al., “Local-first software: you own your data, in spite of the cloud”, 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software, Onward!, Athens, Greece, 2019.

# Adding collaborative features to applications



- replicate the application  $\implies$  dedicated development
- **replicate the application data<sup>a</sup>**
- SQLite is embedded in many applications

---


<sup>a</sup>Kleppmann et al., “Local-first software: you own your data, in spite of the cloud”, 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software, Onward!, Athens, Greece, 2019.

# Referential integrity



- ensure that the **target of a reference exists**
- the deletion of a reference target can result in
  - the **abortion of the deletion**
  - the **propagation of the deletion** to the reference source

# Replicating relational databases: already done?

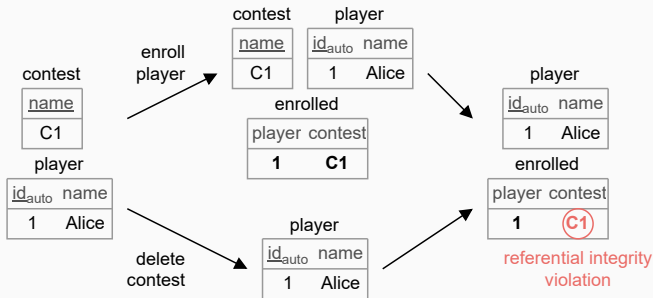


- client-server architecture
- coordination to maintain data integrity<sup>a</sup>

---

<sup>a</sup>Bailis et al., "Highly Available Transactions: Virtues and Limitations", 2013.


# Referential integrity in face of coordination-less concurrencies



- concurrent deletion and referencing of a row



# Coordination-less replication of relational databases

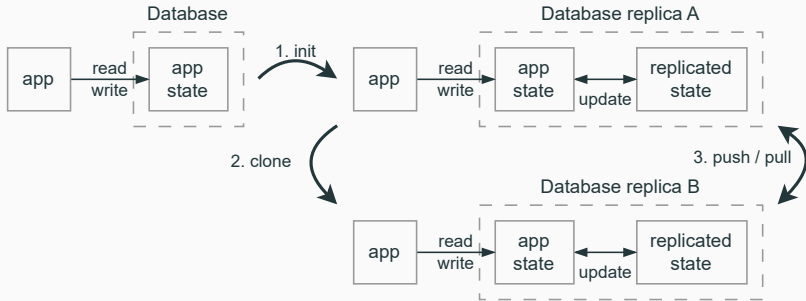


- Git-like **coordination-less** replication of relational databases<sup>a</sup>
- can **break data integrity** and **user intent**
- **not Strongly Convergent**

---

<sup>a</sup>Yu et al., "Conflict-Free Replicated Relations for Multi-Synchronous Database Management at Edge", *IEEE International Conference on Smart Data Services SMDS*, Beijing, China, 2020.


# Coordination-less replication of relational databases



- Git-like **coordination-less** replication of relational databases<sup>a</sup>
- can **break data integrity and user intent**
- **not Strongly Convergent**

<sup>a</sup>Yu et al., "Conflict-Free Replicated Relations for Multi-Synchronous Database Management at Edge", *IEEE International Conference on Smart Data Services SMDS*, Beijing, China, 2020.

# Strong convergence




- property enforced by CRDTs<sup>a</sup>
- advantages:
  - low latency
  - no flickering

---

<sup>a</sup>Shapiro et al., “Conflict-Free Replicated Data Types”, *Stabilization, Safety, and Security of Distributed Systems - 13th International Symposium SSS, Grenoble, France, 2011*.

# Strong convergence

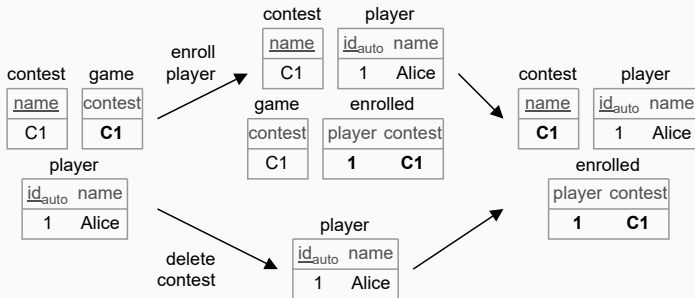


- property enforced by CRDTs<sup>a</sup>
- advantages:
  - low latency
  - no flickering

---

<sup>a</sup>Shapiro et al., “Conflict-Free Replicated Data Types”, *Stabilization, Safety, and Security of Distributed Systems - 13th International Symposium SSS, Grenoble, France, 2011*.

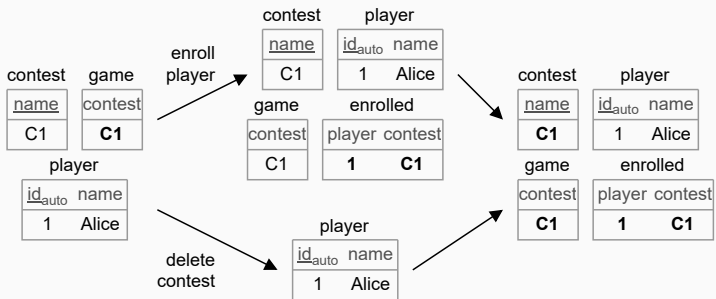
# Referential integrity maintenance - state of the art



- writes are compensated<sup>a</sup> in order to ensure integrity
- the *contest* is restored
- however, the *game* is not restored

<sup>a</sup>Balegas et al., "IPA: Invariant-preserving Applications for Weakly-consistent Replicated Databases", 2018.


# Referential integrity maintenance - desired output



- the *game* should be restored

Can we replicate a relational database without  
any coordination that enforces Strong  
Convergence and maintains data integrity?

# Architecture overview



- app read without overhead
- an app write triggers replicated state update
- push / pull in background
- a pull merges the received state and computes app state



# Replicated state: composing CRDTs



- globally unique and monotonic timestamps
  - monotonic: greater than previously observed timestamps
- Last-Writer-Win (LWW) Register<sup>a</sup> keeps the newest value
- state of CLFlag computed from the longest chain

<sup>a</sup>Johnson et al., "Maintenance of duplicate databases", 1975.

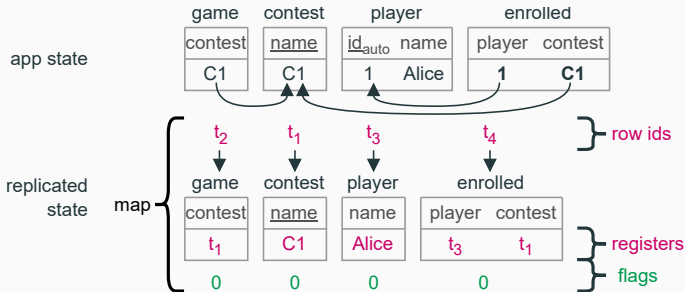
# Replicated state: composing CRDTs



- globally unique and monotonic timestamps
  - monotonic: greater than previously observed timestamps
- Last-Writer-Win (LWW) Register<sup>a</sup> keeps the newest value
- state of CLFlag computed from the longest chain

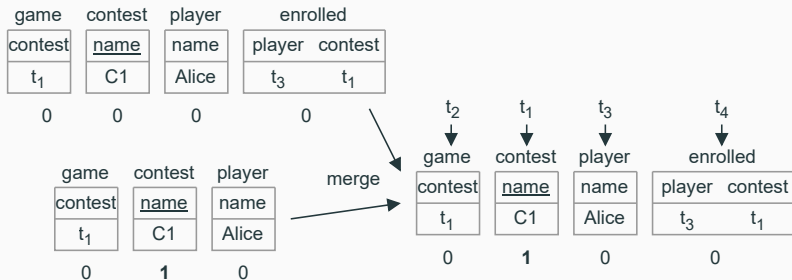
<sup>a</sup>Johnson et al., "Maintenance of duplicate databases", 1975.

# Replicated state: composing CRDTs



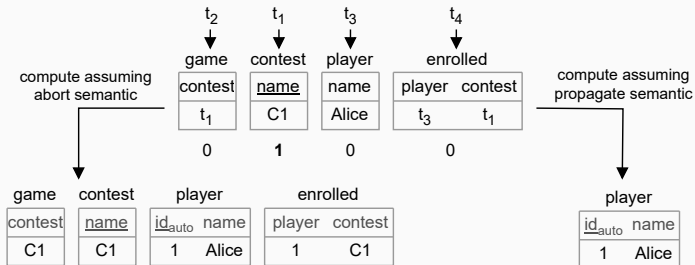
- timestamps as row identifiers
- a CL-Flag indicates if a row is removed
- a replicated attribute is a LWW-Register
- row identifiers as values of foreign keys

# Replicated state: composing CRDTs



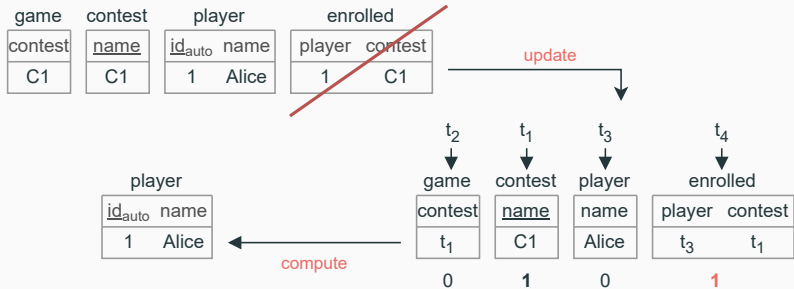
- the replicated state encodes only the app write

# Compute app state from replicated state



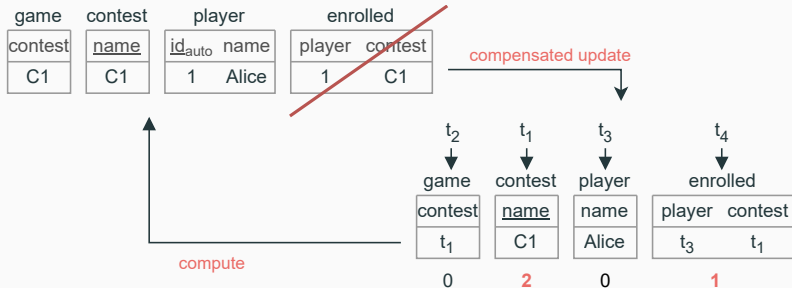
- app state is **derived** from the replicated state
- leverage database schema for selecting **computation semantic**

# Compensation of app writes




- state computation can result in surprising effect on app writes

# Compensation of app writes



- state computation can result in surprising effect on app writes
- **app writes must be compensated** for ensuring user intent

# Conclusions



- **coordination-less** replication of relational database
  - maintains data integrity
  - Strongly Convergent
- composition of CRDTs + state computation + compensations





Victorien Elvinger   4.0